

Mosh

An Interactive Remote Shell for Mobile Clients

Keith Winstein and Hari Balakrishnan

keithw@mit.edu

June 14, 2012

What we built

1. Protocol for low-latency **object synchronization**
 - ▶ with roaming
 - ▶ through suspend/resume
 - ▶ over marginal networks
2. Mobile shell application to replace SSH, telnet.
 - ▶ with “predictive” instant keystroke feedback

Remote terminals

- ▶ 1969: TELNET (RFC 15)
- ▶ 1977: SUPDUP
- ▶ 1991: BSD rlogin
- ▶ 1995: SSH

Secure Shell, 1995

- ▶ Uses TCP.
 - ▶ Connection named by IP:port endpoints.
- ▶ Sends:
 - ▶ user keystrokes → server
 - ▶ octet stream (coded screen updates) → client terminal
- ▶ All UI from server.

Problems with SSH

- ▶ Can't roam:
 - ▶ ... across Wi-Fi networks.
 - ▶ ... from Wi-Fi to cell or vice versa.
- ▶ Times out if data unacknowledged after n minutes.
 - ▶ ... if laptop goes to sleep.
- ▶ Responds poorly to packet loss.

More problems with SSH

- ▶ Byte stream is wrong layer of abstraction.
 - ▶ Client wants *latest* screen.
 - ▶ Don't want to replay megabytes in between.
 - ▶ SSH doesn't understand data, so must send everything.
 - ▶ TCP fills buffers, so Control-C takes forever.
- ▶ Typing and editing on high-latency path is frustrating.
 - ▶ Cellular wireless (100 ms to 500 ms)
 - ▶ Intercontinental (250 ms)
 - ▶ Loaded "4G LTE" (5,000 to 40,000 ms!)


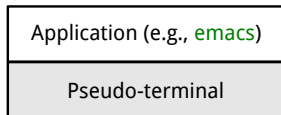
State Synchronization Protocol

- ▶ Runs over UDP.
- ▶ Instead of synchronizing *octet streams*, synchronize *objects*.
- ▶ Object interface:
 - ▶ diff: make vector from state $A \rightarrow B$
 - ▶ patch: apply vector to A to make B
- ▶ Object implementation, **not protocol**, defines synchronization semantics.



State Synchronization Protocol (cont.)

- ▶ Protected by AES-OCB (Krovetz 2011)
 - ▶ Integrity and confidentiality with one key.
- ▶ Key exchange happens out of band.
 - ▶ Uses SSH to bootstrap.
 - ▶ Runs `mosh-server` on remote side.
 - ▶ No privileged code, no daemons.
- ▶ Roaming is easy:
 - ▶ Source address of latest authentic packet from client
⇒ server's new target


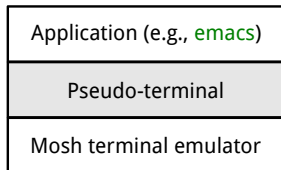
SSH Server






SSH Client



Mosh Server



Mosh Client



Synced
objects



State Synchronization Protocol (cont.)


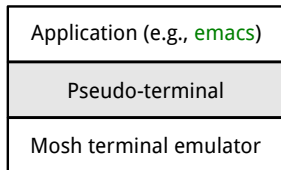
- ▶ **Flow control:** adapt frame rate to network conditions.
- ▶ Minimum interval between frames: smoothed RTT/2.
- ▶ Can skip over states.
- ▶ “P-retransmissions” (post-paper improvement).

P-retransmissions




“Prophylactic” retransmission reduces latency in presence of loss.

1. Last ack was for state #3. Then state changes to #4.
2. Host sends diff from $3 \rightarrow 4$.
3. Object changes to state #5.
4. If no timeout yet, make next diff as $4 \rightarrow 5$.
5. **Also** make diff from $3 \rightarrow 5$: the *p-retransmission*.
6. If retransmission is shorter or not much longer, send it instead.

Mosh Server




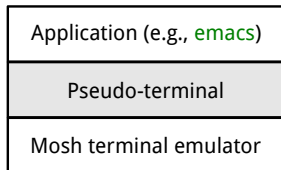
Mosh Client





Synced
objects



Mosh Server



Mosh Client

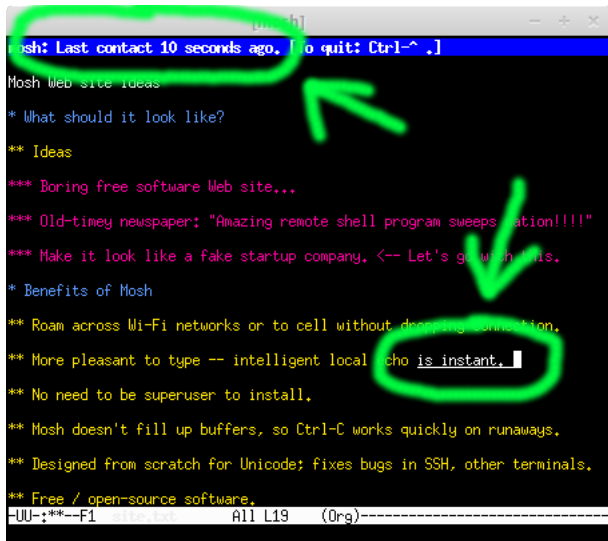


Synced
objects



Speculative Local Echo and Editing

- ▶ Client anticipates server response.
- ▶ Runs predictive model in the background.
- ▶ Make predictions in *epochs*.
- ▶ If any from epoch n is confirmed, show whole epoch.
- ▶ If user does something difficult to handle, become tentative:
increment epoch.
 - ▶ Carriage return
 - ▶ Escape
 - ▶ Up/down arrow
 - ▶ Control char



The screenshot shows a terminal window titled "[mosh]". The first line is a status bar: "mosh: Last contact 10 seconds ago. [no quit: Ctrl-^ .]". This line is highlighted in blue. Below it, the text "Mosh Web site Ideas" is displayed. Then, a prompt "* What should it look like?" is shown. The response "** Ideas" is followed by three lines of suggestions: "*** Boring free software Web site...", "*** Old-timey newspaper: 'Amazing remote shell program sweeps nation!!!!'", and "*** Make it look like a fake startup company. <-- Let's go with this." Next, a prompt "* Benefits of Mosh" is shown. The response "** Roam across Wi-Fi networks or to cell without dropping connection." is followed by "** More pleasant to type -- intelligent local echo is instant.". The text "is instant." is underlined and circled in red. Below this are two more lines: "** No need to be superuser to install." and "** Mosh doesn't fill up buffers, so Ctrl-C works quickly on runaways." Then, "** Designed from scratch for Unicode; fixes bugs in SSH, other terminals." and finally, "** Free / open-source software." The bottom of the window shows a status bar: "-UU-:***--F1 site.txt All L19 (Org)-----". There are two red annotations: a circle around the first status bar and an arrow pointing to the "is instant." line.

```
[mosh]
mosh: Last contact 10 seconds ago. [no quit: Ctrl-^ .]
Mosh Web site Ideas
* What should it look like?
** Ideas
*** Boring free software Web site...
*** Old-timey newspaper: "Amazing remote shell program sweeps nation!!!!"
*** Make it look like a fake startup company. <-- Let's go with this.
* Benefits of Mosh
** Roam across Wi-Fi networks or to cell without dropping connection.
** More pleasant to type -- intelligent local echo is instant.
** No need to be superuser to install.
** Mosh doesn't fill up buffers, so Ctrl-C works quickly on runaways.
** Designed from scratch for Unicode; fixes bugs in SSH, other terminals.
** Free / open-source software.
-UU-:***--F1 site.txt All L19 (Org)-----
```


Demo

Benefits

- ▶ Roaming and suspend/resume:
 - ▶ Sleep and wake up later.
 - ▶ Change networks at will (Wi-Fi, cellular, wired, VPN).
- ▶ Helpful warnings: won't hang without notice.
- ▶ Performance:
 - ▶ Works on marginal links.
 - ▶ Good interactivity even when RTT is > 100 ms.
 - ▶ Semantically appropriate flow control (won't fill up queues, Ctrl-C works quickly, no beeping fits).
- ▶ Security
 - ▶ Uses SSH to bootstrap: no privileged code, no daemons.
- ▶ Correctness
 - ▶ Unicode

Unicode admits varying interpretations.

```
xterm 271
sh$ echo -e "xyz\033[2;2H\0314\0202\nhello"
xyz
hello
sh$
```

```
[mosh]
sh$ echo -e "xyz\033[2;2H\0314\0202\nhello"
xyz
hello
sh$
```

```
GNOME Terminal 3.0.1
sh$ echo -e "xyz\033[2;2H\0314\0202\nhello"
xyz
hello
sh$
```


```
Macintosh HD — Terminal.app 2.2.2
sh$ echo -e "xyz\033[2;2H\0314\0202\nhello"
xyz
ello
sh$
```

bricks the terminal!

Evaluation

- ▶ Collected 40 hours of terminal usage from six users.
- ▶ Covers 10,000 keystrokes using shell, e-mail, text editor (emacs and vi), chat, Web browser.
- ▶ Replayed over:
 1. Sprint 1xEV-DO (3G)
 2. Verizon LTE (4G)
 3. MIT-Singapore
 4. 50% loss path
- ▶ Result: 70% of keystrokes predicted instantly.
- ▶ Prediction errors $< 1\%$

Sprint 1xEV-DO cumulative keystroke response distribution



Evaluation (cont.)

Verizon LTE service in Cambridge, Mass., running one concurrent TCP download:

	Median latency	Mean	σ
SSH	5.36 s	5.03 s	2.14 s
Mosh	< 0.005 s	1.70 s	2.60 s

Deployment

- ▶ Distributed in Debian, Ubuntu, Fedora, Gentoo, Arch, Slackware versions of GNU/Linux.
- ▶ Available via EPEL for Red Hat, CentOS, Oracle Linux.
- ▶ Included in MacPorts, Homebrew, FreeBSD ports collections.
- ▶ Works on Cygwin and Solaris, (very raw) on Android and iOS.
- ▶ News stories in April on Hacker News, Reddit, The Register, Twitter, Slashdot, Barrapunto.
- ▶ Top repository of the month on GitHub.
- ▶ 200,000+ page views, 70,000+ downloads, 1,200+ followers of version control repo.

Reception

@xlfe: “one of those times you don’t realize something is broken until you see it fixed”

@adamhjk: “the user experience really is dreamy.”


@esmolanka: “mosh is awesome. Tested it for two weeks and it really made my life easier: faster feedback and no more reconnects(!)”

@andyd: “Using mosh on the train rather than plain ssh, and it does actually make a huge difference!”

USENIX review: “ISO 2022 locking escape sequences oh flying spaghetti monster please kill me now.”

State Sync Protocol for all?

- ▶ We believe SSP may be appropriate for many network problems.
- ▶ Android Gmail, Google Chat, Skype cannot roam without failure.
- ▶ **June 13, 2012:**



- ▶ Neither can Gmail (Web site).
- ▶ These problems can be expressed as state synchronization.

Next Steps

- ▶ Essay to appear in *;/login:* magazine.
- ▶ Mosh software under development by a team of contributors.
- ▶ We are working to apply SSP to mobile videoconferencing.
- ▶ We hope to show quantitative improvement on standard metrics (latency, quality), plus features like roaming.

Summary

- ▶ SSP is a secure datagram protocol that synchronizes abstract objects across a roaming IP connection.
- ▶ Mosh uses SSP to synchronize a terminal emulator with predictive local echo.
- ▶ In evaluations with 10,000 real-world keystrokes from six users, Mosh markedly reduced user-visible latency across several Internet paths.
- ▶ We think SSP will be useful for other applications as well.
- ▶ <http://mosh.mit.edu>