

# Transport Protocols for Gracefully Mobile Applications

Keith Winstein, Anirudh Sivaraman, and Hari Balakrishnan

M.I.T. CSAIL

November 15, 2012


# Outline

Sprout: Flow control for interactive apps

SSP: Graceful mobility

Alfalfa: video for varying networks

# Mobile wireless networks are variable




# Videoconferencing systems work poorly on LTE

- ▶ We measured cellular networks while driving:
  - ▶ **Verizon LTE**
  - ▶ Verizon 3G (1xEV-DO)
  - ▶ AT&T LTE
  - ▶ T-Mobile 3G (UMTS)
- ▶ Then ran apps across emulated network:
  - ▶ **Skype** (Windows 7)
  - ▶ Google Hangout (Chrome on Windows 7)
  - ▶ Apple Facetime (OS X)

# Why is performance so bad?

- ▶ Exiting schemes **react** to congestion signals.
  - ▶ Packet loss.
  - ▶ Increase in round-trip time.
- ▶ This feedback comes too late to help.
- ▶ The killer: **self-inflicted queueing delay**.
- ▶ Any overshoot means a queue filling up with packets.

# Performance summary



# Sprout's goal


- ▶ As much throughput as possible, with
- ▶ bounded risk of delay  $> 100$  ms.

# Bounded risk of delay

- ▶ **Infer** link speed from interarrival distribution.
- ▶ **Predict** future link speed.
  - ▶ Don't wait for congestion.
- ▶ **Control:** Send as much as possible, but require:
  - ▶ 95% probability all packets will arrive within 100 ms.

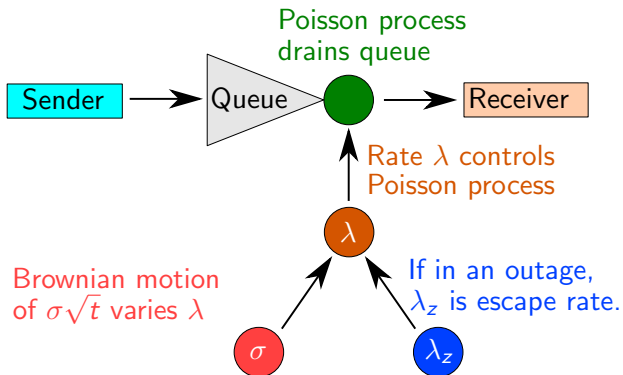


# Infer link speed from interarrival distribution



# Predict future link speed

- ▶ Count packets in every 20 ms tick.
- ▶ Use Bayesian updating to infer (uncertain) link speed.
- ▶ Make a cautious forecast.



# The cautious forecast


- ▶ Receiver has cloud of current link speeds
- ▶ For eight 20 ms ticks in the future:
  - ▶ Predict future link speed
  - ▶ Find 5th percentile of cumulative packets
- ▶ Send forecast to sender (piggyback)
- ▶ Most of the math is precalculated.

# Limitations

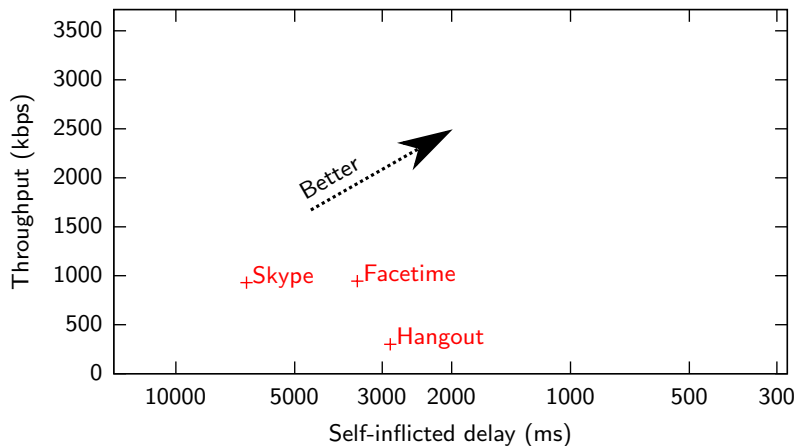
- ▶ Stochastic model has not been tuned
- ▶ Designed for cellular link with per-user queue
- ▶ If other users can cause you big delay, can't solve end-to-end

# Verizon LTE uplink: head-to-head

# Verizon LTE uplink




# Verizon LTE uplink

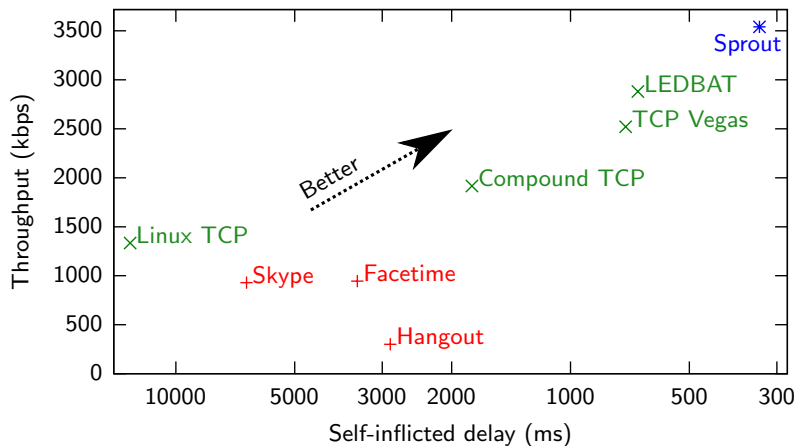




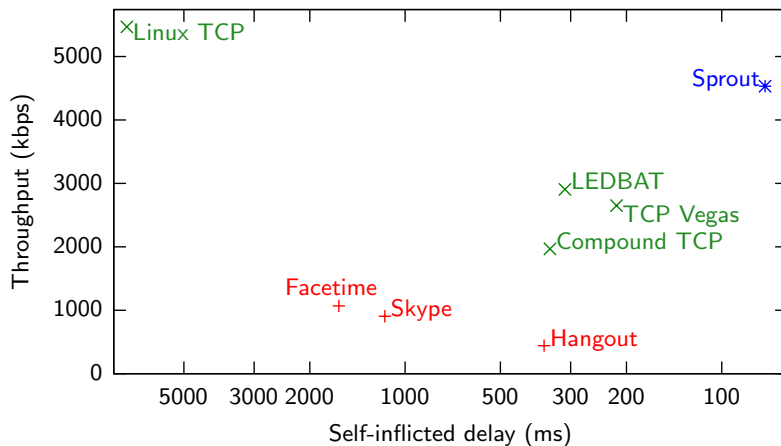
# Verizon LTE uplink



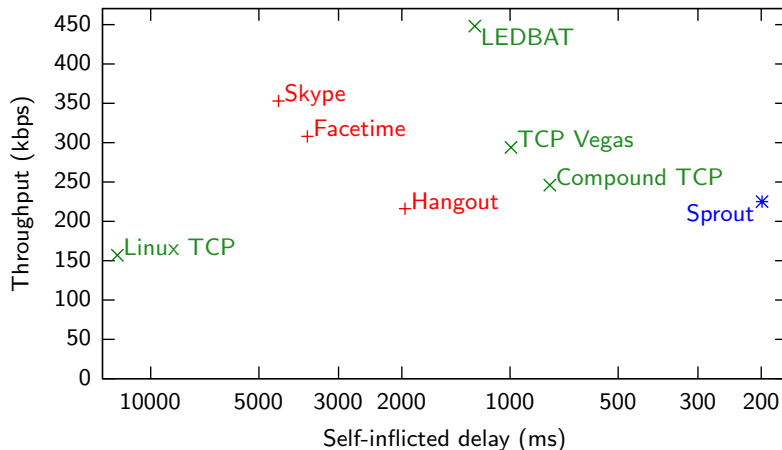
# Verizon LTE uplink




# Verizon LTE downlink




# Verizon 3G (1xEV-DO) uplink



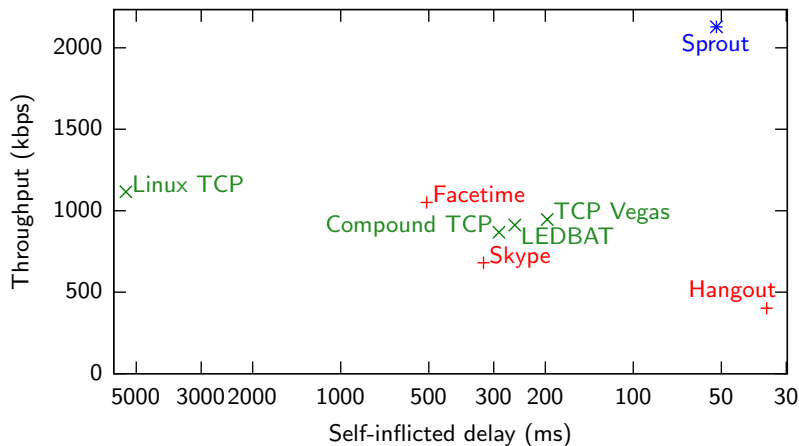
# Verizon 3G (1xEV-DO) downlink



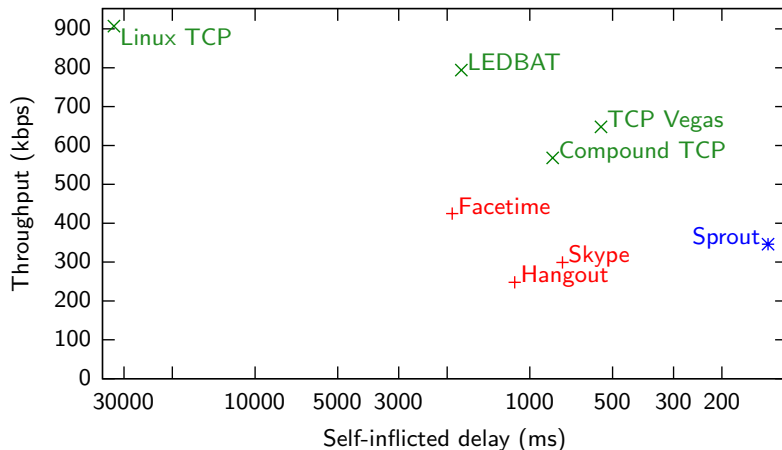
# AT&T LTE uplink



# AT&T LTE downlink

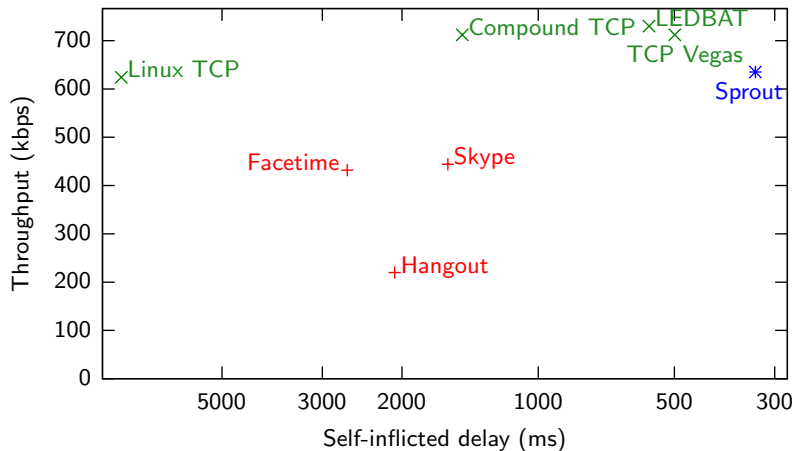


# T-Mobile 3G (UMTS) uplink





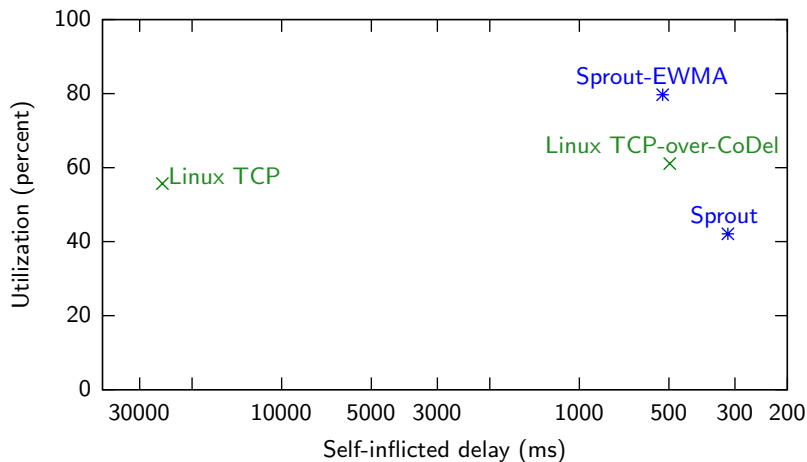
# T-Mobile 3G (UMTS) downlink



# Overall results

Sprout vs.	Avg. speedup	Delay reduction
Skype	2.2×	7.9×
Hangout	4.4×	7.2×
Facetime	1.9×	8.7×
Compound	1.3×	4.8×
TCP Vegas	1.1×	2.1×
LEDBAT	Same	2.8×
Linux TCP (CUBIC)	1.1×	79×

# Competes with AQM even though end-to-end



# Future work

- ▶ Public contest for best predictor
- ▶ Anybody will be able to build protocol using results

# Sprout for controlled delay over cellular networks

- ▶ **Infer** link speed from interarrival distribution
- ▶ **Predict** future link speed
- ▶ **Control** risk of large delay
- ▶ Yields 2–4 $\times$  throughput of Skype, Facetime, Hangout
- ▶ Achieves 7–9 $\times$  reduction in self-inflicted delay
- ▶ Matches active queue management **without router changes**

Sprout: Flow control for interactive apps

SSP: Graceful mobility

Alfalfa: video for varying networks

# Motivation: frustration with SSH

- ▶ Can't roam:
  - ▶ ... across Wi-Fi networks.
  - ▶ ... from Wi-Fi to cell or vice versa.
- ▶ Can't sleep and wake up (usually).
  - ▶ ... TCP disconnects if unacked data for too long.
- ▶ Responds poorly to packet loss.
- ▶ All UI requires round-trip to server.

# Octet stream is wrong layer of abstraction

- ▶ Client wants *latest* screen.
- ▶ After interruption, don't want to replay megabytes.
- ▶ But SSH doesn't understand data, so must send everything.
- ▶ TCP fills buffers, so Control-C takes forever.



# What we built

1. Protocol for low-latency **object synchronization**
  - ▶ with roaming
  - ▶ through suspend/resume
  - ▶ over lossy network paths
2. Supports **rolling latency compensation** on client side
3. Mobile shell application to replace SSH.

# State Synchronization Protocol

- ▶ Runs over UDP.
- ▶ Instead of sending *octet streams*, synchronize *objects*.
- ▶ Object must support:
  - ▶ diff: make vector from state  $A \rightarrow B$
  - ▶ patch: apply vector to  $A$  to make  $B$
- ▶ Object implementation, **not protocol**, defines synchronization semantics.

# Secure quick roaming

- ▶ Protected by AES-OCB (Krovetz 2011)
  - ▶ Integrity and confidentiality with one key.
- ▶ All packets are idempotent operations.
- ▶ Unlike SSH or TLS, connection control is also authenticated.
  - ▶ Attacker cannot terminate connection with FIN or RST.
- ▶ Roaming is easy:
  - ▶ Source address of latest authentic packet from client  
⇒ server's new target
  - ▶ Client may not even **know** it has roamed.

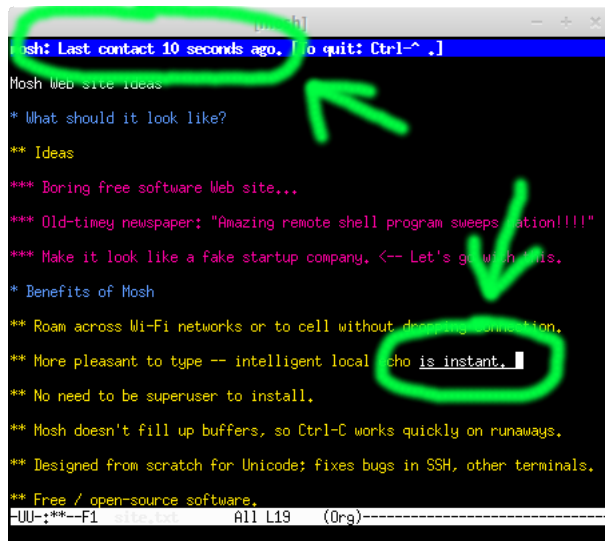
## P·retransmissions trade performance for robustness.

SSP has options in choosing which diff to send:

1. Last ack was for state #3. Then state changes to #4.
2. Host sends diff from  $3 \rightarrow 4$ .
3. Object changes to state #5.
4. If no timeout yet, make next diff as  $4 \rightarrow 5$ .
5. **Also** make diff from  $3 \rightarrow 5$ : the *prophylactic retransmission*.
6. If p·retransmission is shorter or not much longer, send instead.

# Rolling predictions

- ▶ Client runs predictive model of server UI.
- ▶ Make predictions in *epochs*.
- ▶ If any from epoch  $n$  is confirmed, show whole epoch.
- ▶ If prediction from epoch  $n$  is wrong, hide that epoch.
- ▶ If user does something difficult to handle, become tentative:  
*increment epoch*.
- ▶ Better than Meteor's on/off "local mode."



```
mosh: Last contact 10 seconds ago. [to quit: Ctrl-^.]  
Mosh Web site Ideas  
* What should it look like?  
** Ideas  
*** Boring free software Web site...  
*** Old-timey newspaper: "Amazing remote shell program sweeps nation!!!!"  
*** Make it look like a fake startup company. <-- Let's go with this.  
* Benefits of Mosh  
** Roam across Wi-Fi networks or to cell without dropping connection.  
** More pleasant to type -- intelligent local echo is instant.  
** No need to be superuser to install.  
** Mosh doesn't fill up buffers, so Ctrl-C works quickly on runaways.  
** Designed from scratch for Unicode; fixes bugs in SSH, other terminals.  
** Free / open-source software.  
-UU-:***--F1 site.txt All L19 (Org)-----
```

# Demo

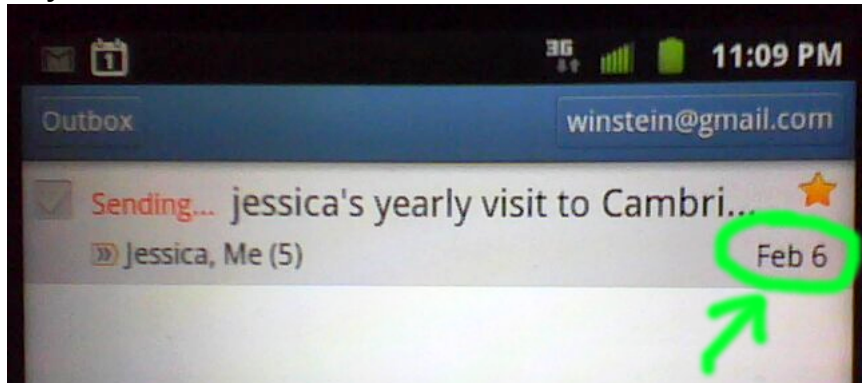
# Deployment

- ▶ In Debian, Ubuntu, Fedora, Gentoo, Arch, Slackware.
- ▶ Available for Red Hat, CentOS, Oracle Linux.
- ▶ In MacPorts, Homebrew, FreeBSD ports collection.
- ▶ Works on Cygwin, Solaris, experimental port to Android.
- ▶ Cover of Linux Magazine this month (Nov. 2012).
- ▶ Top repository of the month on GitHub.
- ▶ 300,000+ page views, 75,000+ downloads, 1,500+ VCS followers.



# Gmail app if user roams at the wrong time

July 5, 2012:



# State synchronization for all

Many Web and native apps have trouble with roaming and intermittent connectivity:

- ▶ Android Gmail app
- ▶ Skype
- ▶ Google Chat
- ▶ gmail.com
- ▶ quora.com
- ▶ Google Voice
- ▶ Twitter

These problems may also be expressed as state synchronization.

Sprout: Flow control for interactive apps

SSP: Graceful mobility

Alfalfa: video for varying networks

# Coded video is also a state synchronization problem

- ▶ P-pictures are “diff” from previously-sent frame.
- ▶ B-pictures are “diff” from two previously-sent frames.

# Mobile video conferencing

For reliable mobile video conferencing, application wants from:

- ▶ **Network:** “How much is it safe to send now?”
  - ▶ TCP and DCCP don’t supply this.
- ▶ **Video encoder:** “Give me a P-picture from this frame to *this* frame with maximum length  $x$ .”
  - ▶ x264 (H.264) and libvpx (VP8) don’t do this.
- ▶ **Video decoder:** “Apply this P-picture to *this* predicate picture, and give me the results.”
  - ▶ QuickTime, libavcodec, etc. don’t allow this.

Solution: Sprout/SSP and “explicit-state” video codec API.

# Video complications

- ▶ P-picture is not just based on predicate *frame*.
- ▶ Also have quantization tables, probability tables, other inter-frame state.
- ▶ Specifications don't envision this use.
- ▶ Aren't explicit about what state needs to be carried across frames.

So far, we have implemented API for MPEG-2 video.

# What about Netflix / YouTube?

Current practice:

- ▶ encode multiple quality levels, each with VBV.
- ▶ player can switch, but requires I-picture and visible jump.
- ▶ often switch only available every 10 seconds!

Our view:

- ▶ VBV was designed for isochronous broadcast channels!
- ▶ Ditch VBV: the player's buffer is all that matters.
- ▶ Encode full trellis of diffs between quality levels.
- ▶ **Let the decoder choose** how it wants to plan ahead.
- ▶ Can run over HTTP.

# Conclusion

- ▶ Sprout: end-to-end flow control for cellular networks that matches or outperforms in-network modifications.
- ▶ SSP: protocol for gracefully mobile state synchronization.
- ▶ Alfalfa: video for varying networks.